

# assgnment\_4\_Nishant

February 3, 2026

```
[7]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
```

```
[11]: df = pd.read_csv("HousingData.csv")
df.head()
```

```
[11]:      CRIM      ZN  INDUS   CHAS      NOX      RM      AGE      DIS      RAD      TAX    PTRATIO \
0  0.00632  18.0    2.31    0.0  0.538  6.575  65.2  4.0900      1    296     15.3
1  0.02731    0.0    7.07    0.0  0.469  6.421  78.9  4.9671      2    242     17.8
2  0.02729    0.0    7.07    0.0  0.469  7.185  61.1  4.9671      2    242     17.8
3  0.03237    0.0    2.18    0.0  0.458  6.998  45.8  6.0622      3    222     18.7
4  0.06905    0.0    2.18    0.0  0.458  7.147  54.2  6.0622      3    222     18.7
```

```
      B    LSTAT    MEDV
0  396.90    4.98  24.0
1  396.90    9.14  21.6
2  392.83    4.03  34.7
3  394.63    2.94  33.4
4  396.90     NaN  36.2
```

```
[13]: features = df.columns.drop("MEDV")
target = "MEDV"
```

```
[15]: df.isnull().sum()
```

```
[15]: CRIM      20
ZN        20
INDUS     20
CHAS      20
NOX       0
RM        0
```

```
AGE      20
DIS      0
RAD      0
TAX      0
PTRATIO   0
B        0
LSTAT    20
MEDV     0
dtype: int64
```

```
[6]: df = df.dropna()
```

```
[7]: df.isnull().sum()
```

```
[7]: CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
MEDV     0
dtype: int64
```

```
[8]: features = df.columns.drop("MEDV")
target = "MEDV"

results = []

for feature in features:
    X = df[[feature]]
    y = df[target]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )

    model = LinearRegression()
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
```

```

r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

results.append([feature, r2, mse])

```

```

[9]: results_df = pd.DataFrame(
    results, columns=["Feature", "R2 Score", "Mean Squared Error"]
)

results_df.sort_values(by="R2 Score", ascending=False)

```

	Feature	R2 Score	Mean Squared Error
5	RM	0.478680	43.971498
12	LSTAT	0.416253	49.236978
10	PTRATIO	0.172987	69.755573
2	INDUS	0.119382	74.276944
1	ZN	0.103697	75.599982
4	NOX	0.075738	77.958155
11	B	0.075087	78.013084
9	TAX	0.046167	80.452434
0	CRIM	0.041634	80.834726
6	AGE	0.003567	84.045536
3	CHAS	-0.023619	86.338587
8	RAD	-0.065704	89.888294
7	DIS	-0.081723	91.239499

```

[17]: import matplotlib.pyplot as plt
import seaborn as sns

target = "MEDV"
features = df.drop(columns=[target]).columns

```

```

[19]: sns.set_theme(style="whitegrid", context="notebook")

for feature in features:
    plt.figure(figsize=(7, 5))

    sns.regplot(
        x=df[feature],
        y=df[target],
        scatter_kws={
            "alpha": 0.6,
            "s": 40,
            "edgecolor": "black"
        },
        line_kws={

```


```

        "color": "darkred",
        "linewidth": 2
    }
)


plt.xlabel(feature, fontsize=12)
plt.ylabel("House Price (MEDV)", fontsize=12)
plt.title(f"{feature} vs House Price", fontsize=14, fontweight="bold")

plt.tight_layout()
plt.show()


```



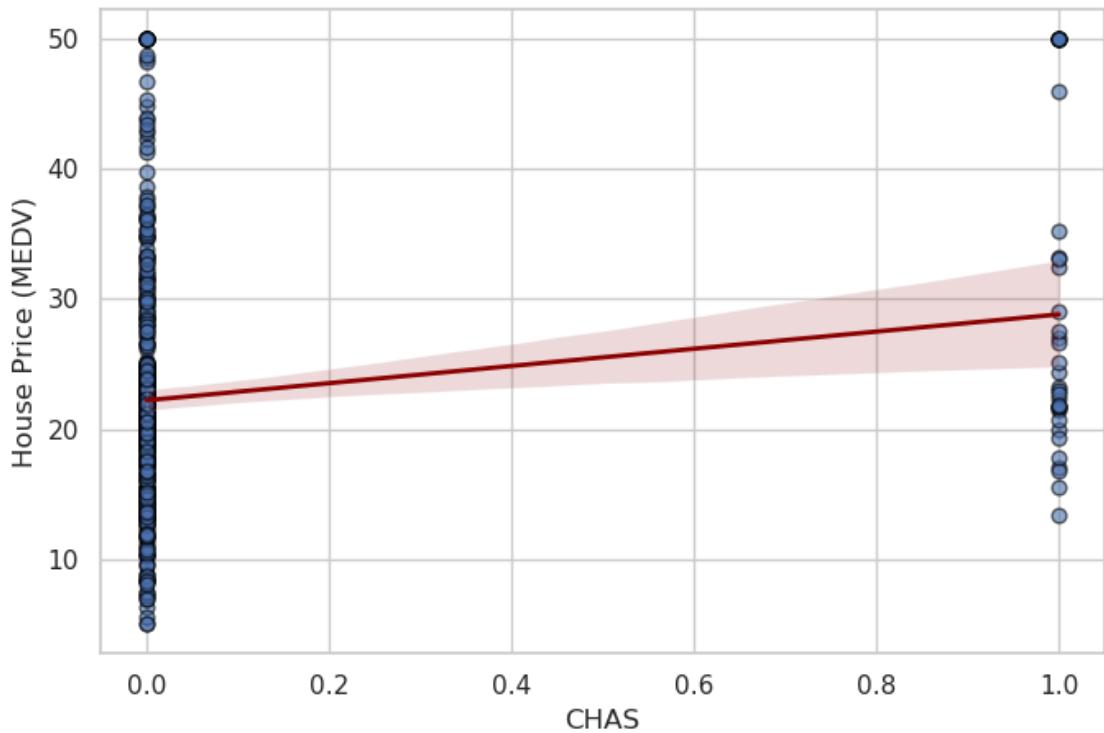
**ZN vs House Price**




**INDUS vs House Price**




**CHAS vs House Price**




### NOX vs House Price




### RM vs House Price




**AGE vs House Price**




**DIS vs House Price**




**RAD vs House Price**




**TAX vs House Price**




**PTRATIO vs House Price**



**B vs House Price**





```
[31]: X = df.drop("MEDV", axis=1)
y = df["MEDV"]
```

```
[45]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )
```

```
[47]: from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy="mean")

X_train = imputer.fit_transform(X_train)
X_test = imputer.transform(X_test)
```

```
[49]: from sklearn.linear_model import LinearRegression

multi_model = LinearRegression()
multi_model.fit(X_train, y_train)
```

```
[49]: LinearRegression()
```

```
[51]: y_pred = multi_model.predict(X_test)

r2_multi = r2_score(y_test, y_pred)
mse_multi = mean_squared_error(y_test, y_pred)

print("R2 Score (Multiple Regression):", r2_multi)
print("Mean Squared Error:", mse_multi)
```

R2 Score (Multiple Regression): 0.659060424186024

Mean Squared Error: 25.00238892351461

```
[15]: importance = pd.DataFrame({
    "Feature": X.columns,
    "Coefficient": multi_model.coef_
})

importance.sort_values(by="Coefficient", ascending=False)
```

```
[15]:   Feature  Coefficient
      5        RM      4.258091
      3       CHAS      1.983837
      8        RAD      0.235588
      1        ZN      0.042440
      2      INDUS      0.025673
     11         B      0.009594
      9       TAX     -0.011997
      6       AGE     -0.021741
      0      CRIM     -0.112187
     12     LSTAT     -0.388620
     10    PTRATIO     -0.975835
      7       DIS     -1.424189
      4      NOX     -17.079257
```

```
[55]: from sklearn.linear_model import LinearRegression

multi_model = LinearRegression()
multi_model.fit(X_train, y_train)
```

```
[55]: LinearRegression()
```

```
[57]: import pandas as pd

importance = pd.DataFrame({
    "Feature": X.columns,
    "Coefficient": multi_model.coef_.flatten()}
```

```
})

[61]: plt.figure(figsize=(11, 7))

importance_sorted = importance.sort_values(by="Coefficient")


# Create a hue column
importance_sorted["Sign"] = importance_sorted["Coefficient"].apply(
    lambda x: "Positive" if x > 0 else "Negative"
)

sns.barplot(
    data=importance_sorted,
    x="Coefficient",
    y="Feature",
    hue="Sign",
    palette={"Positive": "steelblue", "Negative": "indianred"},
    legend=False
)

plt.axvline(0, color="black", linewidth=1)

plt.title("Feature Importance in House Price Prediction", fontsize=16, fontweight="bold")
plt.xlabel("Coefficient Value")
plt.ylabel("Feature")

plt.tight_layout()
plt.show()
```



```
[17]: best_part_a = results_df.sort_values(by="R2 Score", ascending=False).iloc[0]
best_part_a
```

```
[17]: Feature           RM
      R2 Score       0.47868
      Mean Squared Error   43.971498
      Name: 5, dtype: object
```

```
[18]: comparison = pd.DataFrame({
        "Model": ["Best Single Feature (Part A)", "Multiple Features (Part B)"],
        "R2 Score": [best_part_a["R2 Score"], r2_multi],
        "Mean Squared Error": [best_part_a["Mean Squared Error"], mse_multi]
    })

comparison
```

```
[18]:          Model  R2 Score  Mean Squared Error
0  Best Single Feature (Part A)  0.478680      43.971498
1    Multiple Features (Part B)  0.627085      31.454048
```

```
[69]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import pandas as pd
```

```

# Step 1: Convert X_train/X_test back to DataFrame if they were NumPy arrays
X_train = pd.DataFrame(X_train, columns=X.columns)
X_test = pd.DataFrame(X_test, columns=X.columns)

# Step 2: Simple linear regression with one feature, e.g., "RM"
X_simple_train = X_train[["RM"]]
X_simple_test = X_test[["RM"]]

simple_model = LinearRegression()
simple_model.fit(X_simple_train, y_train)

# Step 3: Predict
y_pred_simple = simple_model.predict(X_simple_test)

# Step 4: Compute R2
r2_simple = r2_score(y_test, y_pred_simple)
print("R2 Score (Simple Linear Regression):", r2_simple)

```

R2 Score (Simple Linear Regression): 0.3707569232254778

```
[71]: comparison = pd.DataFrame({
    "Model": ["Simple Linear Regression", "Multiple Regression"],
    "R2 Score": [r2_simple, r2_multi]
})
```

```
[73]: import pandas as pd
import matplotlib.pyplot as plt

# Step 1: Create comparison DataFrame
comparison = pd.DataFrame({
    "Model": ["Simple Linear Regression", "Multiple Regression"],
    "R2 Score": [r2_simple, r2_multi]
})

# Step 2: Assign colors: green for the higher R2, red for lower
colors = ["mediumseagreen" if x == max(comparison["R2 Score"]) else "indianred"
          for x in comparison["R2 Score"]]

# Step 3: Plot
plt.figure(figsize=(6,4))
comparison.set_index("Model")["R2 Score"].plot(
    kind="bar",
    color=colors,
    title="R2 Score Comparison"
)

plt.ylabel("R2 Score")
```

```

plt.xticks(rotation=0)
plt.ylim(0,1) # makes it easier to compare
plt.tight_layout()

# Step 4: Optional: annotate R2 values on top of bars
for i, v in enumerate(comparison["R2 Score"]):
    plt.text(i, v + 0.02, f"{v:.2f}", ha='center', fontweight='bold')

plt.show()

```



```
[21]: best_feature = results_df.sort_values(
    by="R2 Score", ascending=False
).iloc[0]["Feature"]

best_feature
```

```
[21]: 'RM'
```

```
[22]: X_best = df[[best_feature]]
y = df["MEDV"]

X_train_a, X_test_a, y_train_a, y_test_a = train_test_split(
    X_best, y, test_size=0.2, random_state=42
```

```

)
model_a = LinearRegression()
model_a.fit(X_train_a, y_train_a)

y_pred_a = model_a.predict(X_test_a)

```

```

[83]: # Step 1: Compute residuals (Series)
residuals_a = y_test - y_pred_simple          # Single feature model
residuals_b = y_test - y_pred                  # Multiple features model

# Step 2: Plot improved histogram
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,5))

sns.histplot(
    residuals_a,
    bins=25,
    color="tomato",
    label="Part A (Single Feature)",
    kde=True,
    stat="density",
    alpha=0.6
)


sns.histplot(
    residuals_b,
    bins=25,
    color="royalblue",
    label="Part B (Multiple Features)",
    kde=True,
    stat="density",
    alpha=0.6
)

# Add vertical line at 0
plt.axvline(0, color="black", linestyle="--", linewidth=1)

# Labels and title
plt.xlabel("Residuals")
plt.ylabel("Density")
plt.title("Residual Error Comparison: Part A vs Part B", fontsize=14,
          fontweight="bold")
plt.legend()
sns.despine(trim=True)

```

```
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```



[ ]: